

● GTM candidate work sample · 27 Sep 2026

Atlan Agent Registry

Answers to the core challenge

Aryaman Singh · Independent exercise, not affiliated with Atlan

SUMMARY

The position in five lines

ICP

The engineer who owns AI tooling for other people: **platform/DevEx** at most companies, or an **internal AI-enablement engineer** where there's an Enterprise AI team, at software companies of 1,000+ employees that officially allow two or more AI tools.

PRIMARY CHANNEL

Open-source, GitHub-native distribution: a free local CLI (**skillsdrift**) and a public "State of Skill Drift" data drop.

WHY NOW

The big AI tools now govern their own skills, so the pain has moved to the gaps.

HOW ONE BECOMES MANY

Adoption spreads from the engineer to the teams whose skills others depend on. A champion installs the GitHub App, and the Head of AI buys.

VALIDATION

3–5 existing Atlan customers as design partners prove activation. They are where the loop is proven, not where the first 100 come from.

1.4 · WHAT THEY USE TODAY

Governance exists. Inside each tool.


LAYER	TODAY	WHERE IT STOPS
Files	<code>CLAUDE.md</code> <code>AGENTS.md</code> <code>.claude/skills</code> <code>symlinks</code>	Nothing works across repos or tools
AI tool consoles	Claude org skills (review, versions, per-skill usage) · ChatGPT workspace skills · Copilot org instructions · Cursor team rules	Each sees only its own copy
Apps with agents	Gong · Salesforce · Highspot	Copies pasted into assistants are forks
Cloud registries	AWS Agent Registry · Microsoft Agent 365 · Google · Databricks	Built around their own cloud

A year ago the pitch was “you have no governance.” Now each tool governs itself, and **the problem is everything in between.**

SOURCES [Claude org skills](#) [Claude usage analytics](#) [ChatGPT skills](#) [AWS Agent Registry](#) [Microsoft Agent 365](#)

1.5 · WHY THE REGISTRY IS BETTER

The big AI tools now govern their own skills. The pain moved to the gaps.



PROOF · A REAL COMPANY

GitLab already wrote our thesis **down as policy**

"If a leader cannot change the way their function operates by changing a process, they have lost control of their strategy."

GitLab Handbook, Prompts are Process

WHAT THEY HAVE

Claude Enterprise for every team member, plus Glean, Relevance.ai and GitLab Duo

One AI Transformation Owner per function, and a champion community per function

A policy page called "Prompts are Process": every skill must be **shared, versioned, owned**

WHERE IT BREAKS (THEIR WORDS)

The policy is "deliberately tool-agnostic." The tooling is split across 5 platforms.

Shared skills "have to be symlinked into your local skills directory"

"Consume before build" has no catalog spanning every tool

Not a one-off: LinkedIn is hiring "AI Builder, GTM Enablement" roles to "consolidate one-off solutions into durable systems."

SOURCES [Prompts are Process](#) [Hub & Spoke & Hub](#) [AI Transformation Owner](#) [Data team agent setup](#) [LinkedIn AI Builder role](#)

Five triggers turn annoyance into a task this week

01

A second AI tool is approved

Every existing skill needs a second copy, and drift starts.

02

A visible wrong output

Two teammates get different results from “the same” skill.

03

A security finding

Snyk found critical issues in 13.4% of 3,984 skills from ClawHub and skills.sh.

04

An AI operating model launches

An owner per function now exists and needs tooling.

05

A change to how a function works

One change has to reach everyone at once.

THE ONE QUESTION THAT
QUALIFIES A LEAD

“How many AI tools are officially allowed, across how many functions?”

Three people, one job at three altitudes

WHO

IN THEIR WORDS

Engineer

USER

“Keep one version of our AI skills working across every tool and team, and stop being the person everyone asks which one is right.”

Process owner

CHAMPION

“When we change how we work, change it once and have everyone’s AI tool pick it up.”

Head of AI

BUYER

“Show the AI investment works, across every tool we pay for, without losing control of how each function operates.”

Sources: GitLab Handbook · LinkedIn job posts · public developer threads · Atlan’s own demo (“whether that investment is working”)

1.1 · THE FIRST CUSTOMER

The engineer who owns AI tooling **for everyone else**

SOFTWARE COMPANY · 1,000+ EMPLOYEES

≥2 AI TOOLS OFFICIALLY ALLOWED

SHARED AGENT INSTRUCTIONS IN REPOS

**Platform / DevEx engineer, or internal
AI-enablement engineer**

Owens tooling for other people and can change shared repos.

SIGNAL WE CAN BUILD A LIST FROM

CLAUDE.md AGENTS.md .claude/skills

or Cursor/Copilot rule files in ≥3 public repos

or a public job post for AI enablement, an “AI builder”, or an AI Transformation Owner

WHICH WORKLOAD LEADS

Coding

→

Knowledge
work

→

Operational

Coding first (skills live in repos) → knowledge work through hubs → operational agents last.

Not product AI engineers: their problem is production agents, already served by LangSmith and cloud registries.

Who we're **not** targeting yet

NOT YET

WHY

Single-AI-tool companies

Their vendor's console is good enough

Support teams on Decagon / Sierra

The agent is the product, and it's already governed where it runs

Companies under ~1,000 employees

They'll use the free CLI, but there's no second function and no budget owner

Companies all-in on one cloud's agent stack

The native registry wins until they add a second tool


Sellers and marketers as the entry point

No pain, no install rights. They're reached through hubs.

ONE QUESTION QUALIFIES
OR DISQUALIFIES A LEAD

"How many AI tools are officially allowed, **across how many functions?"**

Artifacts travel; people don't.



Individuals don't spread tools inside companies; **hubs do.**

ENGINEERING

1

Local scan

Nothing leaves the machine

2

Team scan

The repos your team owns

3

A forwarded note

"This skill exists in 4 places, in 3 versions"

4

Claim your skills

Naming an owner is the first act of adoption

INTO NON-TECH, WITHOUT A SEPARATE SALE

ACROSS FUNCTIONS

5

Hubs & dependency pull

Data, product marketing, platform, GTM enablement

6

Shared links

A link to the governed version, never a file

7

Champion installs the App

The install confirms a champion already won

→

Head of AI buys

Team inventory and cross-function usage

HOW NON-TECH GETS ONBOARDED

Share a link, not a file. Sharing is onboarding.

01

Drift report

An engineer runs the scan

02

Bridge skill (in Claude / Codex)

Picks skills whose output is used outside engineering. "Who uses what this skill produces?"

03

Slack, with consent

The user approves every message, and owners approve being named

04

Atlan link

Opens the skill's page: versions, owner, usage

05

Sign in with a work account

The recipient is now in Atlan, on the governed version

A file is a copy, and copies are how drift starts. **A link is the governed version.**

Security-flagged skills are never shared. Companies without Atlan get a neutral card, not Atlan branding.

One primary channel: open-source, GitHub-native

WHAT THEY ALREADY DO

Keep skills and agent instructions in git

Discover tools on GitHub, HN and in communities

Run CLIs locally before trusting a service

Forward internal notes to other teams' owners

Add owners through CODEOWNERS

WHAT THE CHANNEL ASKS

Scan them

Find this one there

Run a local, read-only CLI

Forward the drift note

Claim a skill

WHERE THEY FIND TOOLS

Engineers

GitHub · npm · HN · X ·
platform-engineering
communities

Data teams

dbt community (100K+)

Budget owners

LinkedIn · vendors they
already pay · their own
champion

Consumers

They don't look. Tools arrive
via plugins and Slack.

PASSES SECURITY REVIEW

Local

Read-only

No network by default

No telemetry

MIT licence

2.3 · DISCOVERY TO ACTIVATION

Activation means deployment, not sign-up

Discovery

30 seconds,
no sign-up

First value

The report

Team value

Forwarded
notes

Into non- tech

Bridge skill

Champion

The inventory

Decision

Timed by a
trigger

Activation

Manifest
import

A TEAM IS ACTIVATED WHEN

Its skills are in a governed registry, each with a named owner and version; **≥5 people** run them through their AI tools; and that usage can be observed.

“If they stopped tomorrow, would real work break?”

AN ACCOUNT HAS CROSSED OVER AT

≥10

people

≥3

functions

≥2

AI tools

An engineering-only account has stalled where it started.

Five things Atlan builds. Three it partners on.

BUILD (OR CONFIRM IT EXISTS)

- 1 Import the scan manifest.** A scan becomes a workspace in one step (`atlan-registry-import/1` already exists).
- 2 A URL for every skill.** Shared links open it, and whoever opens it gets onboarded.
- 3 Publish into every tool:** Claude, ChatGPT, Codex, Copilot, Cursor, Glean.
- 4 Usage across vendors,** at whatever depth each tool exposes.
- 5 Link skills to data.** Flag a skill when its metric changes. This is the moat.

PARTNER

Scanners (Snyk, Cisco), feeding the approve step

Anthropic and OpenAI directories, as places to publish

Cloud registries (AWS, Google, Microsoft): sync, don't replace

Everything I built works without Atlan. My card links stand in for #2 until Atlan's exist.

2.5 · THE STRONGEST ALTERNATIVES

Strong channels, better used as amplifiers

Public skill directories

e.g. skills.sh

STRONGEST CASE

Real reach: 1.43M all-time installs

WHY NOT PRIMARY

Built for individuals; it optimises the install that causes drift

ROLE INSTEAD

A free amplifier: list the CLI and cards there

Platform partnerships

e.g. AWS

Proves the category at enterprise scale

AWS is building its own registry; partner deals take quarters

Be compatible, not a partner: export in AWS Agent Registry's format

Native agent recommendations

The brief's "can Claude recommend Registry?"

It lives where people already work

It only reaches people already using the tool, so it can't do cold discovery

The loop inside the primary channel.

The bridge skill has Claude recommend and share governed skills with teammates, and that's the step that carries adoption from engineers into business teams.

ALSO WEIGHED

Atlan's existing customers (design partners)

Agencies (later)


Builder communities

2.6 · DOES IT SCALE NON-LINEARLY?

Linear across companies. Non-linear inside each one.

Open question for Atlan: do several teams in one company count toward the 100?

ACROSS COMPANIES



Each first install is a separate decision, spread through content.

INSIDE ONE COMPANY



A team scan covers every repo the team owns, one hub reaches hundreds, and dependency pull brings in the next hub.

RESOURCES

Sources behind every claim

HOW TOOLS GOVERN TODAY

Claude org skills

Claude usage analytics

ChatGPT skills

Custom GPT retirement

REAL-COMPANY PROOF

GitLab: Prompts are Process

GitLab: Hub & Spoke & Hub

GitLab: AI Transformation Owner

LinkedIn: AI Builder, GTM Enablement

SECURITY & REGISTRIES

Snyk ToxicSkills

AWS Agent Registry GA

Microsoft Agent 365

Databricks UC skills

ATLAN & CHANNEL

atlan.com (Agent Skills)

anthropics/skills

skills.sh

dbt community

The full list with claims mapped is in the answers doc appendix (R1-R42). Independent candidate exercise. Accessed Sept 2026.