

Atlan GTM Work Sample: Review

Brief vs. submission evaluation + devil's advocate · Prepared for Aryaman · 27 Sep 2026

Bottom line

Strong on thinking, build, and the AI log. Weak on the brief's third outcome: **shipped evidence**. Right now it reads as a very good strategy doc with a working tool, which is short of what the brief rewards. The enterprise pivot also works against Atlan's stated thesis.

5 strong	6 partial	5 missing
----------	-----------	-----------

1. What was asked vs. what was answered

Ask	What the doc gives	Verdict	Note
Narrowest credible early adopter	Enterprise buying center (4 seats), later widened to non-tech	Missing	Went broadest, not narrowest
Job, trigger, current tools, why better	Triggers (multi-harness drift, key-person, supply-chain) and current tools (copy files, symlinks, sx) solid; 'why better' asserted, not shown	Partial	
User / champion / blocker / budget owner	Four-seat map, clearly defined	Strong	Best part
Individual to team adoption	2-hop chain (engineer to champion to DM) plus share links	Strong	
Who not to target yet	Original exclusion reversed, nothing replaced it	Missing	
Talk to people; use the Atlan team	No customer conversations; memo 401 logged instead of asked for	Missing	Big miss for a 'first week' brief
One channel, beat 2 alternatives	OSS / GitHub-native vs skills.sh and AWS	Strong	
Non-linear scale	Claimed via org App install and share links	Partial	Not demonstrated
Which workload leads (coding / knowledge / ops)	Coding leads, knowledge work expands; ops / CI barely in GTM	Partial	
Campaign 'ready to run, not future tense'	Blocked on X creds, DNS, npm publish, waitlist backend	Missing	Not runnable today
Continue / change / stop signals	Per-hop gates	Partial	Mostly unmeasurable with zero telemetry
Ship evidence beyond a deck (outcome 3)	One public-repo scan stat, zero users	Missing	Weakest area
Something to click / run / inspect	CLI runs locally; no hosted page or waitlist	Partial	
Activation definition + next experiments	Clear, mechanical, smell-tested	Strong	
AI work log	Approach, systems, where AI was wrong, human decisions	Strong	
Understandable in 15 minutes	Not this doc; deck / appendix not reviewed	Partial	Confirm the deck carries it

2. Three things to address before the session

The pivot fights Atlan's own thesis

The brief says: *"We want dramatic distribution before we optimize for a traditional enterprise sales motion."* Correction 1 moves to budget owners, procurement, and a 6-9 month horizon. The two can be reconciled (bottom-up entry, enterprise as destination), but the doc never makes that argument. Expect this as the first pushback.

Numbers that aren't in the brief

"~600-700 orgs, ~20k people, ~30 per org" is not in the challenge PDF, and the brief sets no 30/90-day deadline. If these came from the memo or videos, cite them. If not, the "honest math" section argues against a constraint nobody set.

The brief rewards evidence, and there is none yet

"A rough working system that teaches us something" beats polish. So far the system has taught one thing: the 951-skill scan. Even 3-5 real conversations or 20 real CLI runs would change the read.

Also missed: **native agent recommendations** was on the brief's channel list and never evaluated, even though the #1 skill on skills.sh being a skill-finder is the doc's own strongest pain signal.

3. Devil's advocate on the working doc

1. Atlan's existing customers are ignored

Atlan already sells to MasterCard, GM, and others and has CSMs. The doc weighs HN, skills.sh, and AWS but not its own catalog base, where 'context layer' is a natural cross-sell. Likely the first question.

2. Enterprise ICP vs. HN channel

The target has CISOs; the entry point is an unknown MIT npm CLI from HN. Many enterprises block unvetted packages. 'Passes security review by construction' is not real: review is a process, not a property. The 'not Atlan' disclaimer lowers trust further.

3. The headline stat is weak

'951 skills, 100% ungoverned' is from public libraries, where ownership lives in git / CODEOWNERS, not in the file. The result follows from the definition. HN will point this out, and Day 0 leads with it. 'A 30-person org scans as chaos' is asserted, never measured.

4. The math is circular and the bets are optimistic

~30 people / ~667 orgs were backed out of the target, so they can't 'validate' it. 30% Hop-1 to decision is high: bottom-up to enterprise usually converts in single digits to low teens. Halve both bets and 6-9 months becomes about two years.

5. Zero telemetry makes the gates unmeasurable

'40 orgs touched per wave' can't be counted if nothing phones home. The only signal is a waitlist signup. Tier 2/3 non-tech usage 'needs the product', so activation can't be measured today either.

6. Competition from the platforms is missing

Anthropic, OpenAI, and other harness makers are adding team / admin skill management. Plus git, CODEOWNERS, and plugin marketplaces give a free 'good enough' option. 'The registry is the only cure' is asserted, and platform risk isn't discussed.

7. Hop 1 carries social risk too

The doc says the pincer removes Hop 2's social risk, but Hop 1 asks an engineer to commit a report that says their team's skills have 'no owner, no security review'. People rarely volunteer findings that make their own team look bad.

8. The exec memo overclaims

A local scan can't know whether a security review happened, so 'no security review' is an inference stated as fact. That clashes with the honesty posture. The 11 heuristics also overlap Atlan's announced-but-unshipped scan.

9. The X bot and creator cards could backfire

Unrequested public verdicts on people's repos can read as naming and shaming. A pre-computed card on Joe Karlsson's or the sx author's work may land as a callout, not a gift. X API posting also costs money.

10. The non-tech correction raises the bar early

Requiring 5 eng + 3 GTM + 2 ops and a live non-tech workspace in the pilot's definition of done slows the first 100. The '108x in legal' figures are growth off small bases.

11. It changes the goal it was given

The brief asks for 100 teams. The doc redefines them as enterprise orgs with about 10 deployed people each, then says it can't be done in time. Honest, but can read as 'I can't do the task'. With paid and partnerships ruled out, the only fallback is 'stop and rethink'.

12. There is no demand evidence

Everything is built, nothing is launched: zero scans, zero signups. The biggest open item is whether to launch at all. One week of real data beats 16 rounds of reasoning.

13. The doc is heavy

It uses many coined terms (Hop 1/2, pincer, USP, anti-fragmentation gate, Tier 1/2/3) and spends many words on the process. The answer is buried. Slips: Correction 4 comes before 3, 'hero-sull' typo, 'we' used by a solo candidate.

What holds up

- Buying-center reframe and the four-seat map
- Activation meaning real deployment, with a smell test
- "Artifacts travel, people don't"
- Share links turning the leak into distribution
- The AI verification story, including where the AI was wrong

Fix list, in priority order

1	Launch something small now	Publish the CLI and a waitlist; get real runs before the session.
2	Get 3-5 conversations	Include asking Atlan for the memo, which is what a real first week looks like.
3	Add Atlan's existing customers	Make the catalog base the first channel, or explain why not.
4	Fix or reframe the drift stat	Scan real multi-person setups, or define 'ungoverned' honestly.
5	Reframe enterprise	Enterprise is where the motion ends up, not where it starts. Tie it to 'dramatic distribution first'.
6	Add who not to target	One line is enough.
7	Cut the core doc by half	It needs to read in 15 minutes; move process detail to the appendix.